

# ExploreDuplicates

April 6, 2025

This file is an exploration of duplicate values that are seen in the Iceberg Tracking Beacon Database.

Derek Mueller

Several beacons have data with repeat positions but with consecutive dates, which seem suspicious. For example, the repeats seem to appear in patterns. In my experience, there is almost always some jitter in the GPS data and a lot of jitter in ARGOS data. If they are erroneous, is there a way to reprocess these? See for example:

- 2017\_300234062328750 SVP-I-BXGSA-L-AD
- 2016\_300234063515450 iCALIB
- 2009\_300034012571050 ICEB-I-XA

Here is a bit of the 300234062328750 SVP data. Note that the distance and direction rounding was turned off to generate this.

datetime_data	latitude	longitude	temperature	airdistance	speed	direction
2017-07-25 18:00:00	76.3194	-75.0602	5.4	1607.1913	0.446	92.36
2017-07-25 19:00:00	76.3194	-75.0602	7.0	0	0.0	180
2017-07-25 20:00:00	76.3194	-75.0602	9.8	0	0.0	180
2017-07-25 21:00:00	76.295	-74.993	9.8	3251.8739	0.903	146.86
2017-07-25 22:00:00	76.295	-74.993	9.3	0	0.0	180
2017-07-25 23:00:00	76.295	-74.993	8.8	0	0.0	180
2017-07-26 00:00:00	76.2778	-74.9708	9.7	2007.9831	0.557	162.97
2017-07-26 01:00:00	76.2778	-74.9708	8.9	0	0.0	180
2017-07-26 02:00:00	76.2778	-74.9708	5.4	0	0	180

The following notebook will review duplicates in the ITDB.

```
[1]: # imports
```

```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import pyproj
```

```
[2]: def count_decimal_places(value):
      if "." in str(value): # Check if there is a decimal point
```

```

    return len(
        str(value).split(".")[1]
    ) # Count the characters after the decimal point
return 0

```

```

[3]: def lon_precision_v_distance(lat):
    """
    Generate distances represented by a 1 sd change in longitude.

    Parameters
    -----
    lat : float
        A valid latitude

    Returns
    -----
    dist :

    """
    geodesic = pyproj.Geod(ellps="WGS84")

    lons = [
        -90.1,
        -90.2,
        -80.01,
        -80.02,
        -70.001,
        -70.002,
        -60.0001,
        -60.0002,
        -50.00001,
        -50.00002,
        -40.000001,
        -40.000002,
        -30.0000001,
        -30.0000002,
    ]

    decimal_places = pd.Series(range(1, int(len(lons) / 2) + 1))
    az, baz, dist = geodesic.inv(
        [np.nan] + lons[:-1],
        [lat] * len(lons),
        lons,
        [lat] * len(lons),
    )
    dist = pd.Series(dist)
    distance_m = dist[dist < dist.quantile(0.51)].reset_index(drop=True)

```

```

    return pd.DataFrame({"decimal_places": decimal_places, "distance_m":
↳distance_m})

```

```

[ ]: ### MAIN

# Disable rounding to see the distance properly.
# export the database to csv and read it in.
df = pd.read_csv("/ibtd/20250406/20250406.csv")

print(f"The database has {len(df)} iceberg positions")

```

```

[12]: # make a duplicate indicator
df["dup"] = 0
df.loc[(df["speed"] == 0) & (df["direction"].isna()), "dup"] = 1

```

```

[13]: print( f"{df.dup.sum()} or {df.dup.sum()/len(df):.8%} of these positions are
↳duplicates, where there is no apparent movement of the iceberg"
)
# Apply function to the column to get decimal places
df["lat_d"] = df["latitude"].apply(count_decimal_places)
df["lon_d"] = df["longitude"].apply(count_decimal_places)

```

74590 or 8.97922824% of these positions are duplicates, where there is no apparent movement of the iceberg

The main question is:

Are the duplicates there because there is actually no movement or are there artifacts in the data?

The precision of the data makes a big difference, when detecting duplicates. For more info, see <https://xkcd.com/2170/>, but realize that the latitude has a big effect. For example:

```

[14]: print(
    f"At 80 deg N, the longitude decimal place affects distance as follows:
↳\n\n {lon_precision_v_distance(80)}"
)
print(
    f"At 70 deg N, the longitude decimal place affects distance as follows:
↳\n\n {lon_precision_v_distance(70)}"
)
print(
    f"At 60 deg N, the longitude decimal place affects distance as follows:
↳\n\n {lon_precision_v_distance(60)}"
)
print(
    f"At 50 deg N, the longitude decimal place affects distance as follows:
↳\n\n {lon_precision_v_distance(50)}"
)

```

)

At 80 deg N, the longitude decimal place affects distance as follows:

	decimal_places	distance_m
0	1	1939.348314
1	2	193.934855
2	3	19.393486
3	4	1.939349
4	5	0.193935
5	6	0.019393
6	7	0.001939

At 70 deg N, the longitude decimal place affects distance as follows:

	decimal_places	distance_m
0	1	3818.653700
1	2	381.865412
2	3	38.186541
3	4	3.818654
4	5	0.381865
5	6	0.038187
6	7	0.003819

At 60 deg N, the longitude decimal place affects distance as follows:

	decimal_places	distance_m
0	1	5579.999626
1	2	558.000015
2	3	55.800002
3	4	5.580000
4	5	0.558000
5	6	0.055800
6	7	0.005580

At 50 deg N, the longitude decimal place affects distance as follows:

	decimal_places	distance_m
0	1	7169.574828
1	2	716.957536
2	3	71.695754
3	4	7.169575
4	5	0.716958
5	6	0.071696
6	7	0.007170

Given the location precision of a single frequency (L1) GPS receiver is typically +/- 3 to 10 m, it is quite possible that movement below ~15 m would not be detectable and therefore we cannot expect to separate duplicates that are legitimate from those that are caused by artifacts at SD <= 4.

The number of recorded latitudes by precision (number of decimal places)

```
[16]: df.groupby("lat_d").size()
```

```
[16]: lat_d
1      4878
2     34013
3    278696
4    198913
5     81738
6    155024
7       847
8     35715
9         1
10     1250
13    30450
14     9170
dtype: int64
```

The number of recorded longitudes by precision (number of decimal places)

```
[17]: df.groupby("lon_d").size()
```

```
[17]: lon_d
1      1817
2     30847
3    268560
4    215490
5     90707
6    155337
7       923
8     35346
9         1
10     1253
13    30388
14        7
15       19
dtype: int64
```

The percent of records that are duplicated by precision

```
[18]: df.loc[df["dup"] == 1].groupby("lat_d").size() / df.groupby("lat_d").size() * 100
```

```
[18]: lat_d
1      4.202542
2     5.156852
3     7.282846
4    22.851196
5     6.268810
```

```

6      0.382521
7      NaN
8      3.189136
9      NaN
10     NaN
13     NaN
14     0.261723
dtype: float64

```

```

[19]: df.loc[df["dup"] == 1].groupby("lon_d").size() / df.groupby("lon_d").size() * 100

```

```

[19]: lon_d
1      2.036324
2     10.150096
3      3.367962
4     25.744582
5      5.816530
6      0.364369
7      NaN
8      2.990437
9      NaN
10     0.079808
13     NaN
14     NaN
15     NaN
dtype: float64

```

Next figure out which beacon tracks have the most duplicates:

```

[24]: # find the total number of positions by track
total_counts = df.groupby("beacon_id").size().reset_index()
total_counts.columns = ["beacon_id", "total_n"]

# find the number of duplicates by track
dup_counts = df.loc[df.dup == 1].groupby("beacon_id").size().reset_index()
dup_counts.columns = ["beacon_id", "dups_n"]

# get the beacon models
mf = pd.read_excel("/ibtd/metadata/track_metadata_raw.ods")
mf = mf[["beacon_id", "model"]]

# get the median precision for the duplicates in the track
dp_median = df.loc[df.dup == 1].groupby("beacon_id").agg({'lon_d':
    'median', 'lat_d': 'median'}).reset_index()
#df_median = df.groupby("beacon_id").agg({'lon_d': 'median', 'lat_d': 'median'}).
    reset_index()

```

```
#dp_dup_stats = dup.groupby("beacon_id").agg({"lon_d": ["min", "mean", "max"],
↳ "lat_d": ["min", "mean", "max"]}).reset_index()
#dp_df_stats = df.groupby("beacon_id").agg({"lon_d": ["min", "mean", "max"],
↳ "lat_d": ["min", "mean", "max"]}).reset_index()

# merge dfs to make one with duplicate counts, total counts, beacon model and
↳ percent
totdf = pd.merge(mf, total_counts, how="left")
dupdf = pd.merge(totdf, dup_counts, how="left")
# get the % of track that has duplicates
dupdf["dups_percent"] = dupdf["dups_n"] / dupdf["total_n"] * 100
# merge with the median precision
dupdp = pd.merge(dupdf, dp_median, how="left")
# sort
dupdp.sort_values("dups_percent", inplace=True, ascending=False)
```

```
[25]: n= pd.get_option('display.max_rows')
pd.set_option('display.max_rows', len(dupdp.loc[dupdp["dups_percent"] > 5]))
dupdp.loc[dupdp["dups_percent"] > 5]
#pd.set_option('display.max_rows', n)
```

```
[25]:
```

	beacon_id	model	total_n	dups_n	\
115	2018_300234066545280	FT-2000	6058	4938.0	
109	2017_300234063516450	iCALIB	2429	1840.0	
70	2015_300234061762030	iCALIB	3427	2298.0	
62	2014_300234060544160	SVP-I-BXGS-LP	3107	1923.0	
106	2017_300234062325760	SVP-I-BXGSA-L-AD	461	282.0	
108	2017_300234062328750	SVP-I-BXGSA-L-AD	3165	1923.0	
68	2015_300234060104820	SVP-I-BXGS-LP	8533	4688.0	
223	2023_300534064036660	iCALIB	8220	4504.0	
63	2014_300234061763040	iCALIB	2654	1271.0	
38	2011_300234010035940b	SVP-I-XXGS-LP	31900	14723.0	
105	2017_300234062324750	SVP-I-BXGSA-L-AD	481	221.0	
44	2012_1000000000000000	DMR-800L	603	248.0	
89	2016_300234063513450	iCALIB	11638	4669.0	
42	2011_300234010958690b	SVP-I-XXGS-LP	8124	2977.0	
90	2016_300234063515450	iCALIB	12554	4238.0	
69	2015_300234060435010	SVP-I-BXGS-LP	2311	716.0	
86	2016_300234062950220	SVP-I-BXGS-LP	13676	4142.0	
88	2016_300234062957250	SVP-I-BXGS-LP	10882	3066.0	
21	2010_300034013723350	Model 703 Ice Tracking Buoy	456	126.0	
20	2010_300034013721340	Model 703 Ice Tracking Buoy	67	13.0	
107	2017_300234062327750	SVP-I-BXGSA-L-AD	3117	565.0	
85	2016_300234061768060	iCALIB	2874	517.0	
51	2012_300234010082470	SVP-I-XXGS-LP	2342	328.0	
15	2009_300034012571050	ICEB-I-XA	3858	422.0	

203	2021_300234011751690	iCALIB	18022	1731.0
207	2021_300234060725890	iCALIB	11516	1100.0
204	2021_300234011751700	iCALIB	12752	1128.0
202	2021_300234011750710	iCALIB	12650	1046.0
201	2021_300234011750690	iCALIB	12947	1020.0
34	2011_300034013463170	Model 703 Ice Tracking Buoy	5433	384.0
87	2016_300234062951220	SVP-I-BXGS-LP	12275	774.0
22	2010_300034013726340	Model 703 Ice Tracking Buoy	524	32.0
56	2013_300034013464170	Model 703 Ice Tracking Buoy	4922	286.0
55	2013_300034013464160	Model 703 Ice Tracking Buoy	7647	433.0
205	2021_300234011752700	iCALIB	13718	775.0
200	2019_300434063495310	ITB v2.0	2909	161.0
199	2019_300434063494100	ITB v2.0	5171	281.0

	dups_percent	lon_d	lat_d
115	81.512050	5.0	5.0
109	75.751338	4.0	4.0
70	67.055734	4.0	4.0
62	61.892501	4.0	4.0
106	61.171367	4.0	3.0
108	60.758294	4.0	4.0
68	54.939646	4.0	4.0
223	54.793187	4.0	4.0
63	47.889977	4.0	4.0
38	46.153605	4.0	4.0
105	45.945946	3.0	4.0
44	41.127695	4.0	4.0
89	40.118577	4.0	3.0
42	36.644510	4.0	4.0
90	33.758165	4.0	4.0
69	30.982259	4.0	4.0
86	30.286634	4.0	3.0
88	28.174968	4.0	4.0
21	27.631579	4.0	4.0
20	19.402985	4.0	4.0
107	18.126404	4.0	4.0
85	17.988866	4.0	4.0
51	14.005124	4.0	4.0
15	10.938310	5.0	5.0
203	9.604927	4.0	4.0
207	9.551928	4.0	4.0
204	8.845671	4.0	4.0
202	8.268775	4.0	4.0
201	7.878273	4.0	4.0
34	7.067918	4.0	4.0
87	6.305499	4.0	4.0
22	6.106870	4.0	4.0



56	5.810646	8.0	8.0
55	5.662351	8.0	8.0
205	5.649512	4.0	4.0
200	5.534548	6.0	6.0
199	5.434152	6.0	6.0

## 1 Conclusions

The GPSDELAY has now been accounted for, this has allayed concerns a lot. Clearly there were many metocean beacons that had reporting frequencies that were much less than their transmission interval. See below for the head of the table above.

Note that before GPSDELAY was addressed the table was like this (first few rows). Before at least 23 tracks had duplicates  $> 10\%$ .

beacon_id	beacon_model	total_n	dups_n	dups_percent
2015_300234061762030	iCALIB	10280	9150	89.01
2014_300234060544160	SVP-I-BXGS-LP	9324	8139	87.29
2017_300234062325760	SVP-I-BXGSA-L-AD	1381	1202	87.04
2017_300234062328750	SVP-I-BXGSA-L-AD	9499	8255	86.90
2015_300234060104820	SVP-I-BXGS-LP	25682	21836	85.02
2014_300234061763040	iCALIB	7961	6577	82.62
2017_300234062324750	SVP-I-BXGSA-L-AD	1442	1182	81.97

After reviewing the current top offenders above (which do not show ‘jerky’ stop start motion, it seems that the duplicates here are due to non-motion (which is the way it is supposed to be).

Closing the case on this.

[ ]: